

PRACTICA-ELASTICSEARCH-HADOOP

Santiago Cisneros

KeepCoding

Parte 1:

The screenshot shows the Google Cloud console interface for a Dataproc cluster named 'practica-arquitectura'. The cluster is in the 'En ejecución' (Running) state. The 'INSTANCIAS' tab is selected, showing three instances: 'practica-arquitectura-m', 'practica-arquitectura-w-0', and 'practica-arquitectura-w-1'. To the right, an SSH terminal window is open, displaying the Ubuntu 20.04.6 LTS welcome message and system information. The terminal output shows the user 'santiago_keeppoding' logging in and running the command 'gsutil cp gs://bucket-practica-elasticsearch/elasticsearch-hadoop-8.14.1.jar .'. The command is successful, and the file is copied to the local directory.

Parte 2:

The screenshot shows the Google Cloud console interface for a Dataproc cluster named 'practica-elastica'. The cluster is in the 'En ejecución' (Running) state. The 'INSTANCIAS' tab is selected, showing three instances: 'practica-elastica-m', 'practica-elastica-w-0', and 'practica-elastica-w-1'. To the right, an SSH terminal window is open, displaying the Ubuntu 20.04.6 LTS welcome message and system information. The terminal output shows the user 'santiago_keeppoding' logging in and running the command 'gsutil cp gs://bucket-practica-elasticsearch/elasticsearch-hadoop-8.14.1.jar .'. The command is successful, and the file is copied to the local directory. The terminal also shows the output of the 'curl' command, which returns a 200 OK status and the content type 'application/json'.

Parte 3:

The screenshot shows the Google Cloud Console interface for the 'bd14-project'. The 'Compute Engine' section is active, displaying a list of VM instances. The instances listed are 'practica-arquitectura-m', 'practica-arquitectura-w-0', 'practica-arquitectura-w-1', and 'practica-elastic'. The 'practica-elastic' instance is highlighted. To the right, an SSH terminal window is open, showing the execution of various commands to configure the Elasticsearch service, including copying files, setting environment variables, and starting the service.

Google Cloud Console - Compute Engine

Estado de prueba gratuita: crédito por \$275.65 y 77 días restantes. Activa tu cuenta completa para obtener acceso a recursos y paga solo por lo que usas.

bd14-project

Compute Engine

Instancias de VM

INSTANCIAS OBSERVABILIDAD PROGRAMAS DE T...

Instancias de VM

Filtro Ingresar el nombre o el valor de la propiedad

Estado	Nombre	Zona
✓	practica-arquitectura-m	us-central1-a
✓	practica-arquitectura-w-0	us-central1-a
✓	practica-arquitectura-w-1	us-central1-a
✓	practica-elastic	us-east1-c

Acciones relacionadas

- Explorar Backup and DR (NUEVO)
- Consulta el informe de facturación
- Supervisa VMs

SSH en el navegador

```
santiago_keeppoding@practica-arquitectura-m:~$ gsutil cp gs://bucket-practica-elasticsearch/commons-httpclient-3.1.jar .
Copying gs://bucket-practica-elasticsearch/commons-httpclient-3.1.jar...
[1 files][297.8 KiB/297.8 KiB]
Operation completed over 1 objects/297.8 KiB.
santiago_keeppoding@practica-arquitectura-m:~$ curl -I http://34.138.133.4:9200
HTTP/1.1 200 OK
X-elastic-product: Elasticsearch
content-type: application/json
content-length: 541

santiago_keeppoding@practica-arquitectura-m:~$ sudo sed -i '5d' /etc/hive/conf.dist/hive-site.xml
santiago_keeppoding@practica-arquitectura-m:~$ sudo sed -i '5a \ <property>\n <name>es.nodes</name>\n <value>34.138.133.4</value>\n </property>\n' /etc/hive/conf.dist/hive-site.xml
santiago_keeppoding@practica-arquitectura-m:~$ sudo sed -i '5a \ <property>\n <name>es.port</name>\n <value>9200</value>\n </property>\n' /etc/hive/conf.dist/hive-site.xml
santiago_keeppoding@practica-arquitectura-m:~$ sudo sed -i '5a \ <property>\n <name>es.nodes.wan.only</name>\n <value>true</value>\n </property>\n' /etc/hive/conf.dist/hive-site.xml
santiago_keeppoding@practica-arquitectura-m:~$ sudo sed -i '5a \ <property>\n <name>hive.aux.jars.path</name>\n <value>/usr/lib/hive/lib/elasticsearch-hadoop-0.14.1.jar:/usr/lib/hive/lib/commons-httpclient-3.1.jar</value>\n </property>\n' /etc/hive/conf.dist/hive-site.xml
santiago_keeppoding@practica-arquitectura-m:~$ sudo cp elasticsearch-hadoop-0.14.1.jar /usr/lib/hive/lib/
santiago_keeppoding@practica-arquitectura-m:~$ sudo cp commons-httpclient-3.1.jar /usr/lib/hive/lib/
santiago_keeppoding@practica-arquitectura-m:~$ sudo service hive-server2 restart
santiago_keeppoding@practica-arquitectura-m:~$
```

Parte 4:

The screenshot shows the Google Cloud Console interface for the 'bd14-project'. The 'Compute Engine' section is active, displaying a list of VM instances. The 'practica-elastic' instance is highlighted. To the right, an SSH terminal window is open, showing the execution of various commands to configure the Elasticsearch service, including copying files, setting environment variables, and starting the service.

SSH en el navegador

```
verification_mode: certificate
keystore.path: certs/transport-p12
truststore.path: certs/transport-p12
# Create a new cluster with the current node only
# Additional nodes can still join the cluster later
cluster.initial_master_nodes: ["practica-elastic"]

# Allow HTTP API connections from anywhere
# Connections are encrypted and require user authentication
http.host: 0.0.0.0

# Allow other nodes to join the cluster from anywhere
# Connections are encrypted and mutually authenticated
#transport.host: 0.0.0.0

#----- END SECURITY AUTO CONFIGURATION -----
santiago_keeppoding@practica-elastic:~$ curl -X POST "localhost:9200/_doc/6" -H 'Content-Type: application/json' -d '{
  "title": "New Document",
  "content": "This is a new document for the master class",
  "tag": ["general", "testing"]
}'
{"_index": "alumnos", "_id": "6", "_version": 1, "result": "created", "_shards": {"total": 2, "successful": 1, "failed": 0}, "_score": 1.0}
santiago_keeppoding@practica-elastic:~$
```

SSH en el navegador

```
{
  "index": {
    "_index": "alumnos",
    "_id": "5"
  },
  "_source": {
    "name": "Luis",
    "last_name": "Martinez"
  },
  "_type": "alumnos",
  "_id": "5",
  "_version": 1,
  "result": "created",
  "_shards": {"total": 2, "successful": 1, "failed": 0},
  "_seq_no": 2,
  "primary_term": 1,
  "status": 201
},
{
  "index": {
    "_index": "alumnos",
    "_id": "7"
  },
  "_source": {
    "name": "Sofia",
    "last_name": "Ramirez"
  },
  "_type": "alumnos",
  "_id": "7",
  "_version": 1,
  "result": "created",
  "_shards": {"total": 2, "successful": 1, "failed": 0},
  "_seq_no": 3,
  "primary_term": 1,
  "status": 201
},
{
  "index": {
    "_index": "alumnos",
    "_id": "8"
  },
  "_source": {
    "name": "Pedro",
    "last_name": "Hernández"
  },
  "_type": "alumnos",
  "_id": "8",
  "_version": 1,
  "result": "created",
  "_shards": {"total": 2, "successful": 1, "failed": 0},
  "_seq_no": 4,
  "primary_term": 1,
  "status": 201
},
{
  "index": {
    "_index": "alumnos",
    "_id": "6"
  },
  "_source": {
    "title": "New Document",
    "content": "This is a new document for the master class",
    "tag": ["general", "testing"]
  },
  "_type": "alumnos",
  "_id": "6",
  "_version": 1,
  "result": "created",
  "_shards": {"total": 2, "successful": 1, "failed": 0},
  "_seq_no": 5,
  "primary_term": 1,
  "status": 201
}
}
```

SSH en el navegador

SUBIR ARCHIVO

DESCARGAR ARCHIVO

```
verification mode: certificate
keystore.path: certs/transport.p12
truststore.path: certs/transport.p12
# Create a new cluster with the current node only
# Additional nodes can still join the cluster later
cluster.initial_master_nodes: ["practica-elastic"]

# Allow HTTP API connections from anywhere
# Connections are encrypted and require user authentication
http.host: 0.0.0.0

# Allow other nodes to join the cluster from anywhere
# Connections are encrypted and mutually authenticated
#transport.host: 0.0.0.0

#----- END SECURITY AUTO CONFIGURATION -----
santiago_keeppoding@practica-elastic:~$ curl -X POST "localhost:9200/alumnos/_doc/" -H 'Content-Type: application/json' -d '{
  > "title": "New Document",
  > "content": "This is a new document for the master class",
  > "tag": ["general", "testing"]
  > }'
{"_index": "alumnos", "_id": "6", "version": 1, "result": "created", "_shards": {"total": 2, "successful": 1, "failed": 0}, "_type": "application/json"}
santiago_keeppoding@practica-elastic:~$
```

practica-elastic

us-east1-c

Migrate to Virtual Machin...

Almacenamiento

Marketplace

Notas de versión

Acciones relacionadas

Explorar Backup and DR

Consul facturación

Transfiere archivos a una VM de

SSH en el navegador

SUBIR ARCHIVO

DESCARGAR ARCHIVO

```
santiago_keeppoding@practica-elastic:~$ curl -X POST "localhost:9200/alumnos/_doc/" -H 'Content-Type: application/json' -d '{
  > "title": "New Document",
  > "content": "This is a new document for the master class",
  > "tag": ["general", "testing"]
  > }'
{"_index": "alumnos", "_id": "6", "version": 1, "result": "created", "_shards": {"total": 2, "successful": 1, "failed": 0}, "_type": "application/json"}
santiago_keeppoding@practica-elastic:~$
```

practica-elastic

us-east1-c

Migrate to Virtual Machin...

Almacenamiento

Marketplace

Notas de versión

Acciones relacionadas

Explorar Backup and DR

Consul facturación

Transfiere archivos a una VM de

SSH en el navegador

SUBIR ARCHIVO

DESCARGAR ARCHIVO

```
santiago_keeppoding@practica-architecture-m:~$ curl -X POST "localhost:9200/alumnos/_doc/" -H 'Content-Type: application/json' -d '{
  > "title": "New Document",
  > "content": "This is a new document for the master class",
  > "tag": ["general", "testing"]
  > }'
{"_index": "alumnos", "_id": "6", "version": 1, "result": "created", "_shards": {"total": 2, "successful": 1, "failed": 0}, "_type": "application/json"}
santiago_keeppoding@practica-architecture-m:~$
```

practica-architecture-m

us-east1-c

Migrate to Virtual Machin...

Almacenamiento

Marketplace

Notas de versión

Acciones relacionadas

Explorar Backup and DR

Consul facturación

Transfiere archivos a una VM de

Parte 5:

